

Delhivery Business Case Study

Introduction

Delhivery, India's leading and rapidly growing integrated player, has set its sights on creating the commerce operating system. They achieve this by utilizing world-class infrastructure, ensuring the highest quality in logistics operations, and harnessing cutting-edge engineering and technology capabilities.

What is expected

The company wants to understand and process the data coming out of data engineering pipelines: \ • Clean, sanitize and manipulate data to get useful features out of raw fields \ • Make sense out of the raw data and help the data science team to build forecasting models on it.

1. Data

The analysis was done on the data located at -

https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/001/551/original/delhivery_data.csv?1642751181

2. Libraries

Below are the libraries required for analysing and visualizing data

```
In [1]: # libraries to analyze data
import numpy as np
import pandas as pd
import scipy.stats as sps

# libraries to visualize data
import matplotlib.pyplot as plt
import seaborn as sns

# Misc libraries
import random
```

3. Data loading and exploratory data analysis

Loading the data into Pandas dataframe for easily handling of data

```
In [2]: # read the file into a pandas dataframe
df = pd.read_csv('delhivery_data.csv')
# look at the datatypes of the columns
print('*****')
print(df.info())
print('*****\n')
print('*****')
```

```

print(f'Shape of the dataset is {df.shape}')
print('*****\n')
print('*****')
print(f'Number of nan/null values in each column: \n{df.isna().sum()}')
print('*****\n')
print('*****')
print(f'Number of unique values in each column: \n{df.nunique()}')
print('*****\n')
print('*****')
print(f'Duplicate entries: \n{df.duplicated().value_counts()}')
print('*****')

```

```

*****
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144867 entries, 0 to 144866
Data columns (total 24 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   data                                     144867 non-null  object
1   trip_creation_time                     144867 non-null  object
2   route_schedule_uuid                   144867 non-null  object
3   route_type                             144867 non-null  object
4   trip_uuid                             144867 non-null  object
5   source_center                         144867 non-null  object
6   source_name                           144574 non-null  object
7   destination_center                    144867 non-null  object
8   destination_name                      144606 non-null  object
9   od_start_time                        144867 non-null  object
10  od_end_time                          144867 non-null  object
11  start_scan_to_end_scan                144867 non-null  float64
12  is_cutoff                            144867 non-null  bool
13  cutoff_factor                        144867 non-null  int64
14  cutoff_timestamp                     144867 non-null  object
15  actual_distance_to_destination        144867 non-null  float64
16  actual_time                          144867 non-null  float64
17  osrm_time                           144867 non-null  float64
18  osrm_distance                       144867 non-null  float64
19  factor                              144867 non-null  float64
20  segment_actual_time                  144867 non-null  float64
21  segment_osrm_time                   144867 non-null  float64
22  segment_osrm_distance                144867 non-null  float64
23  segment_factor                      144867 non-null  float64
dtypes: bool(1), float64(10), int64(1), object(12)
memory usage: 25.6+ MB
None

```

```

*****
Shape of the dataset is (144867, 24)
*****

```

```

*****
Number of nan/null values in each column:
data                                     0
trip_creation_time                     0
route_schedule_uuid                   0
route_type                             0
trip_uuid                             0
source_center                         0
source_name                           293
destination_center                    0
destination_name                      261
od_start_time                        0
od_end_time                          0
start_scan_to_end_scan                0
is_cutoff                            0

```

```

cutoff_factor          0
cutoff_timestamp        0
actual_distance_to_destination  0
actual_time             0
osrm_time               0
osrm_distance           0
factor                 0
segment_actual_time     0
segment_osrm_time       0
segment_osrm_distance    0
segment_factor          0
dtype: int64
*****

```

```
*****
```

```

Number of unique values in each column:
data                2
trip_creation_time  14817
route_schedule_uuid 1504
route_type          2
trip_uuid           14817
source_center       1508
source_name         1498
destination_center  1481
destination_name    1468
od_start_time       26369
od_end_time         26369
start_scan_to_end_scan 1915
is_cutoff           2
cutoff_factor       501
cutoff_timestamp    93180
actual_distance_to_destination 144515
actual_time         3182
osrm_time           1531
osrm_distance       138046
factor              45641
segment_actual_time  747
segment_osrm_time    214
segment_osrm_distance 113799
segment_factor       5675
dtype: int64
*****

```

```
*****
```

```

Duplicate entries:
False      144867
Name: count, dtype: int64
*****

```

```
In [3]: # look at the top 5 rows
df.head()
```

```
Out[3]:
```

| | data | trip_creation_time | route_schedule_uuid | route_type | trip_uuid | source_center | source_name |
|---|----------|-------------------------------|---|------------|--------------------|---------------|-------------|
| 0 | training | 2018-09-20 02:35:36.476840 | thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3... | Carting | 153741093647649320 | IND388121AAA | Anand_VUN |
| 1 | training | 2018-09-20 02:35:36.476840 | thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3... | Carting | 153741093647649320 | IND388121AAA | Anand_VUN |
| 2 | training | 2018-09-20 02:35:36.476840 | thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3... | Carting | 153741093647649320 | IND388121AAA | Anand_VUN |

| | | | | | | | |
|---|----------|-------------------------------|---|---------|-----------------------------|--------------|-----------|
| 3 | training | 2018-09-20 02:35:36.476840 | thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3... | Carting | trip- 153741093647649320 | IND388121AAA | Anand_VUN |
| 4 | training | 2018-09-20 02:35:36.476840 | thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3... | Carting | trip- 153741093647649320 | IND388121AAA | Anand_VUN |

5 rows × 24 columns

Insight

- A quick look at the information of the data reveals that there are **144867 rows and 24 columns** implying 144867 trips have been made with each trip having information such as *trip_creation_time*, *trip_uuid*, *source_center*, *source_name*, *destination_center*, *destination_name* to name a few. Most of the datatype are either "object" or "float64" except for *is_cutoff* and *cutoff_factor*.
- We can also infer that **there are 293 missing values or null value in source_name and 261 missing values or null value in destination_name** in the dataset. As these numbers are small compared to dataset size, 144867, it is safe to drop the rows with the missing values
- There are **no duplicate entries**.
- As columns *is_cutoff*, *cutoff_factor*, *cutoff_timestamp*, *factor* and *segment_factor* are Unknown fields, there is no harm in dropping these columns.
- It makes sense to convert columns *data* and *route_type* to "category" datatype
- It makes sense to convert columns *trip_creation_time*, *od_start_time*, *od_end_time* to "datetime" datatype

```
In [4]: df = df.dropna(how='any')
df = df.drop(columns = ["is_cutoff", "cutoff_factor", "cutoff_timestamp", "factor", "seg
df["data"] = df["data"].astype("category")
df["route_type"] = df["route_type"].astype("category")
df["trip_creation_time"] = pd.to_datetime(df["trip_creation_time"], format='%Y-%m-%d %H:
df["od_start_time"] = pd.to_datetime(df["od_start_time"], format='%Y-%m-%d %H:%M:%S.%f')
df["od_end_time"] = pd.to_datetime(df["od_end_time"], format='%Y-%m-%d %H:%M:%S.%f')
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 144316 entries, 0 to 144866
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   data                                  144316 non-null category
1   trip_creation_time                   144316 non-null datetime64[ns]
2   route_schedule_uuid                 144316 non-null object
3   route_type                          144316 non-null category
4   trip_uuid                           144316 non-null object
5   source_center                       144316 non-null object
6   source_name                         144316 non-null object
7   destination_center                  144316 non-null object
8   destination_name                    144316 non-null object
9   od_start_time                       144316 non-null datetime64[ns]
10  od_end_time                         144316 non-null datetime64[ns]
11  start_scan_to_end_scan              144316 non-null float64
12  actual_distance_to_destination      144316 non-null float64
13  actual_time                         144316 non-null float64
14  osrm_time                          144316 non-null float64
15  osrm_distance                      144316 non-null float64
16  segment_actual_time                 144316 non-null float64
17  segment_osrm_time                  144316 non-null float64
```

```
18 segment_osrm_distance      144316 non-null float64
dtypes: category(2), datetime64[ns](3), float64(8), object(6)
memory usage: 20.1+ MB
```

```
In [6]: df.describe()
```

```
Out[6]:
```

| | trip_creation_time | od_start_time | od_end_time | start_scan_to_end_scan | actual_distance_to_destination |
|-------|-------------------------------|-------------------------------|-------------------------------|------------------------|--------------------------------|
| count | 144316 | 144316 | 144316 | 144316.000000 | 144316.000000 |
| mean | 2018-09-22 13:05:09.454117120 | 2018-09-22 17:32:42.435769344 | 2018-09-23 09:36:54.057172224 | 963.697698 | 234.708 |
| min | 2018-09-12 00:00:16.535741 | 2018-09-12 00:00:16.535741 | 2018-09-12 00:50:10.814399 | 20.000000 | 9.000 |
| 25% | 2018-09-17 02:46:11.004421120 | 2018-09-17 07:37:35.014584832 | 2018-09-18 01:29:56.978912 | 161.000000 | 23.352 |
| 50% | 2018-09-22 03:36:19.186585088 | 2018-09-22 07:35:23.038482944 | 2018-09-23 02:49:00.936600064 | 451.000000 | 66.135 |
| 75% | 2018-09-27 17:53:19.027942912 | 2018-09-27 22:01:30.861209088 | 2018-09-28 12:13:41.675546112 | 1645.000000 | 286.919 |
| max | 2018-10-03 23:59:42.701692 | 2018-10-06 04:27:23.392375 | 2018-10-08 03:00:24.353479 | 7898.000000 | 1927.447 |
| std | NaN | NaN | NaN | 1038.082976 | 345.480 |

Insight

- The data is provided from **2018-09-12 00:00:16.535741 to 2018-10-03 23:59:42.701692**
- The **average time** taken to deliver from source to destination is **964 mins** with **least time being 20mins** and **maximum time being 7898 mins**
- The **average distance** between source and destination warehouse is **235 Kms** with **least distance being 9 Kms** and **maximum distance being 1927 Kms**

4. Detailed Analysis

4.1. Detecting outliers

4.1.1. Outliers for every continuous variable

```
In [7]: # helper function to detect outliers
def detectOutliers(df):
    q1 = df.quantile(0.25)
    q3 = df.quantile(0.75)
    iqr = q3-q1
    lower_outliers = df[df<(q1-1.5*iqr)]
    higher_outliers = df[df>(q3+1.5*iqr)]
    return lower_outliers, higher_outliers
```

```
In [8]: numerical_columns = ['start_scan_to_end_scan', 'actual_distance_to_destination', 'actual_distance_to_destination']
column_outlier_dictionary = {}
for column in numerical_columns:
    print('*'*50)
    print(f'Outliers of \'{column}\'' column are:')
```

```

lower_outliers, higher_outliers = detectOutliers(df[column])
print("Lower outliers:\n", lower_outliers)
print("Higher outliers:\n", higher_outliers)
print('*'*50, end="\n")
column_outlier_dictionary[column] = [lower_outliers, higher_outliers]

```

Outliers of 'start_scan_to_end_scan' column are:

Lower outliers:

Series([], Name: start_scan_to_end_scan, dtype: float64)

Higher outliers:

```

32950      3897.0
32951      3897.0
32952      3897.0
32953      3897.0
32954      3897.0

```

...

```

79524      4239.0
79525      4239.0
79526      4239.0
79527      4239.0
123196      7898.0

```

Name: start_scan_to_end_scan, Length: 373, dtype: float64

Outliers of 'actual_distance_to_destination' column are:

Lower outliers:

Series([], Name: actual_distance_to_destination, dtype: float64)

Higher outliers:

```

402      704.090688
403      726.181078
404      748.332196
405      770.365887
406      796.335857

```

...

```

144796      1611.171536
144797      1633.419313
144798      1650.202066
144799      1673.310381
144800      1689.639499

```

Name: actual_distance_to_destination, Length: 17818, dtype: float64

Outliers of 'actual_time' column are:

Lower outliers:

Series([], Name: actual_time, dtype: float64)

Higher outliers:

```

407      1241.0
408      1277.0
409      1305.0
410      1322.0
411      1352.0

```

...

```

144796      2640.0
144797      2675.0
144798      2700.0
144799      2736.0
144800      2784.0

```

Name: actual_time, Length: 16507, dtype: float64

Outliers of 'osrm_time' column are:

Lower outliers:

Series([], Name: osrm_time, dtype: float64)

Higher outliers:

```

405      630.0

```

```
406          641.0
407          655.0
408          671.0
409          696.0
...
144796      1492.0
144797      1512.0
144798      1532.0
144799      1549.0
144800      1508.0
Name: osrm_time, Length: 17406, dtype: float64
*****
*****
Outliers of 'osrm_distance' column are:
Lower outliers:
  Series([], Name: osrm_distance, dtype: float64)
Higher outliers:
  405          850.4080
  406          865.7213
  407          886.1183
  408          908.4596
  409          944.6344
...
144796      1980.0975
144797      2008.9586
144798      2036.3992
144799      2059.0195
144800      2063.7663
Name: osrm_distance, Length: 17547, dtype: float64
*****
*****
Outliers of 'segment_actual_time' column are:
Lower outliers:
  1805         -26.0
  3761         -21.0
  39825        -58.0
  40942       -211.0
  56464        -12.0
  58697        -36.0
  70479        -42.0
  73603        -51.0
  85042       -244.0
  100205       -74.0
  119377       -48.0
  125821       -16.0
  142409       -15.0
Name: segment_actual_time, dtype: float64
Higher outliers:
  21           93.0
  34           94.0
  72           75.0
  73           78.0
  106          79.0
...
144790         83.0
144819         88.0
144848        302.0
144853         91.0
144866        268.0
Name: segment_actual_time, Length: 9249, dtype: float64
*****
*****
Outliers of 'segment_osrm_time' column are:
Lower outliers:
  Series([], Name: segment_osrm_time, dtype: float64)
Higher outliers:
```

```


34          70.0
38          45.0
157         81.0
158         81.0
214         44.0
...
144802      48.0
144829      74.0
144837      42.0
144843      43.0
144845      54.0
Name: segment_osrm_time, Length: 6348, dtype: float64
*****
*****
Outliers of 'segment_osrm_distance' column are:
Lower outliers:
Series([], Name: segment_osrm_distance, dtype: float64)
Higher outliers:
34          72.5561
157         79.6653
158         82.4127
214         52.7136
316         60.0755
...
144774      60.6393
144802      61.0445
144829      70.0436
144837      60.4795
144845      55.6993
Name: segment_osrm_distance, Length: 4295, dtype: float64
*****

```

```

In [9]: df[numerical_columns].boxplot(figsize=(25,12))
plt.show()

```



```

In [10]: for key, value in column_outlier_dictionary.items():
print(f'The column \'{key}\'' has {len(value[0]) + len(value[1])} outliers')

```

```

The column 'start_scan_to_end_scan' has 373 outliers
The column 'actual_distance_to_destination' has 17818 outliers
The column 'actual_time' has 16507 outliers
The column 'osrm_time' has 17406 outliers
The column 'osrm_distance' has 17547 outliers
The column 'segment_actual_time' has 9262 outliers

```


The column 'segment_osrm_time' has 6348 outliers
The column 'segment_osrm_distance' has 4295 outliers

Insight

- I will not be removing any outliers now.

4.1.2. Remove the outliers

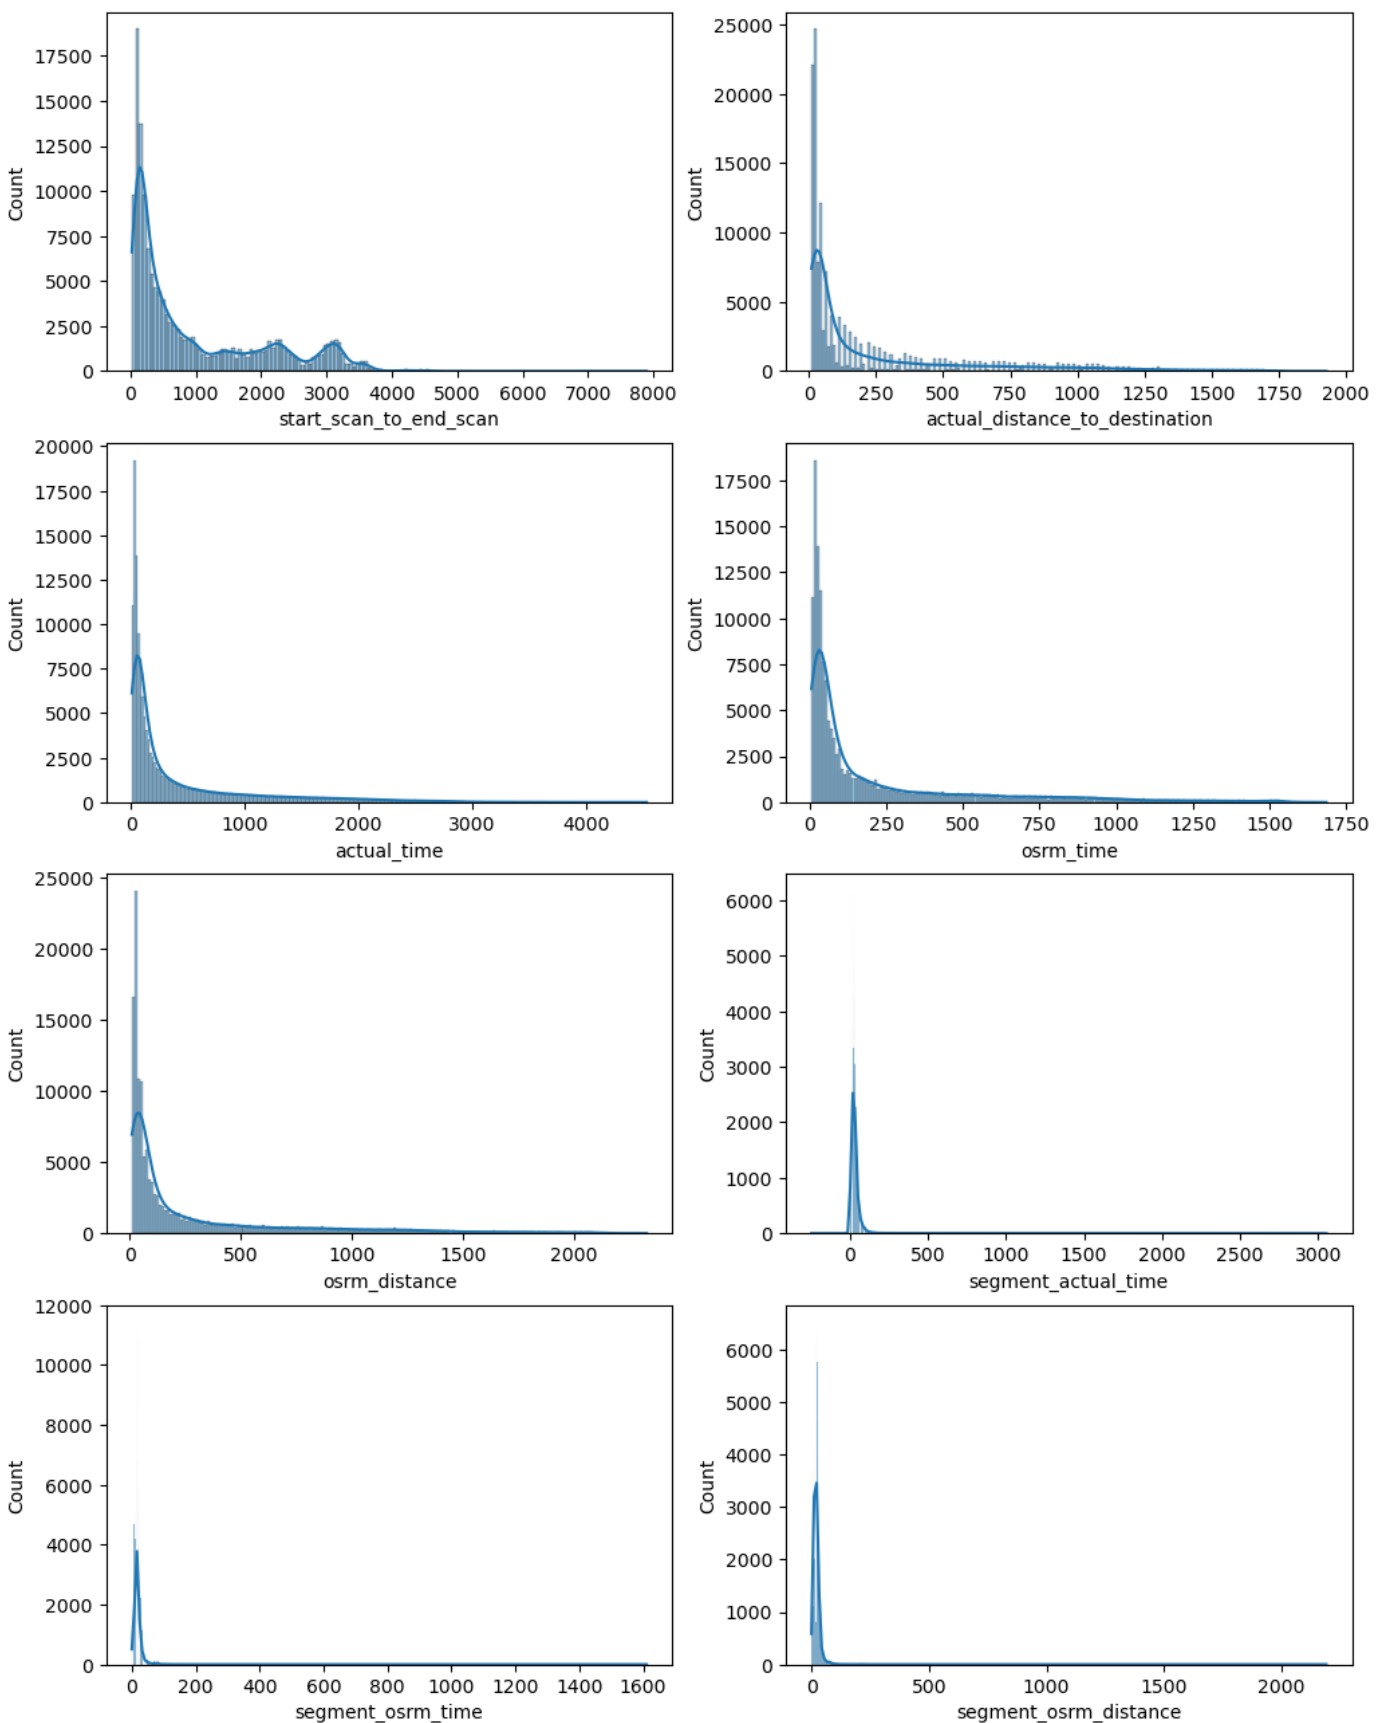
```
In [11]: remove_outliers = False
if True == remove_outliers:
    for key, value in column_outlier_dictionary.items():
        lower_outliers = value[0]
        higher_outliers = value[1]
        df.drop(lower_outliers.index, inplace=True)
        df.drop(higher_outliers.index, inplace=True)
else:
    print('Not removing any outliers')
```

Not removing any outliers

4.2. Univariate analysis

4.2.1. Numerical Variables


```
In [12]: fig, ax = plt.subplots(nrows=4, ncols=2, figsize = (12, 16))
sns.histplot(data=df, x = "start_scan_to_end_scan", kde=True, ax=ax[0,0])
sns.histplot(data=df, x = "actual_distance_to_destination", kde=True, ax=ax[0,1])
sns.histplot(data=df, x = "actual_time", kde=True, ax=ax[1,0])
sns.histplot(data=df, x = "osrm_time", kde=True, ax=ax[1,1])
sns.histplot(data=df, x = "osrm_distance", kde=True, ax=ax[2,0])
sns.histplot(data=df, x = "segment_actual_time", kde=True, ax=ax[2,1])
sns.histplot(data=df, x = "segment_osrm_time", kde=True, ax=ax[3,0])
sns.histplot(data=df, x = "segment_osrm_distance", kde=True, ax=ax[3,1])
plt.show()
```



4.2.2. Categorical Variables

```
In [13]: categorical_columns = ["data", "route_type"]
plt.figure(figsize=(6,6))
plt.subplot(2,1,1)
data = df["data"].value_counts()
plt.pie(data.values, labels = data.index, autopct='%.1f%%')
plt.title("data")
plt.subplot(2,1,2)
```

```
data = df["route_type"].value_counts()
plt.pie(data.values, labels = data.index, autopct='%0.1f%%')
plt.title("route_type")
plt.show()
```

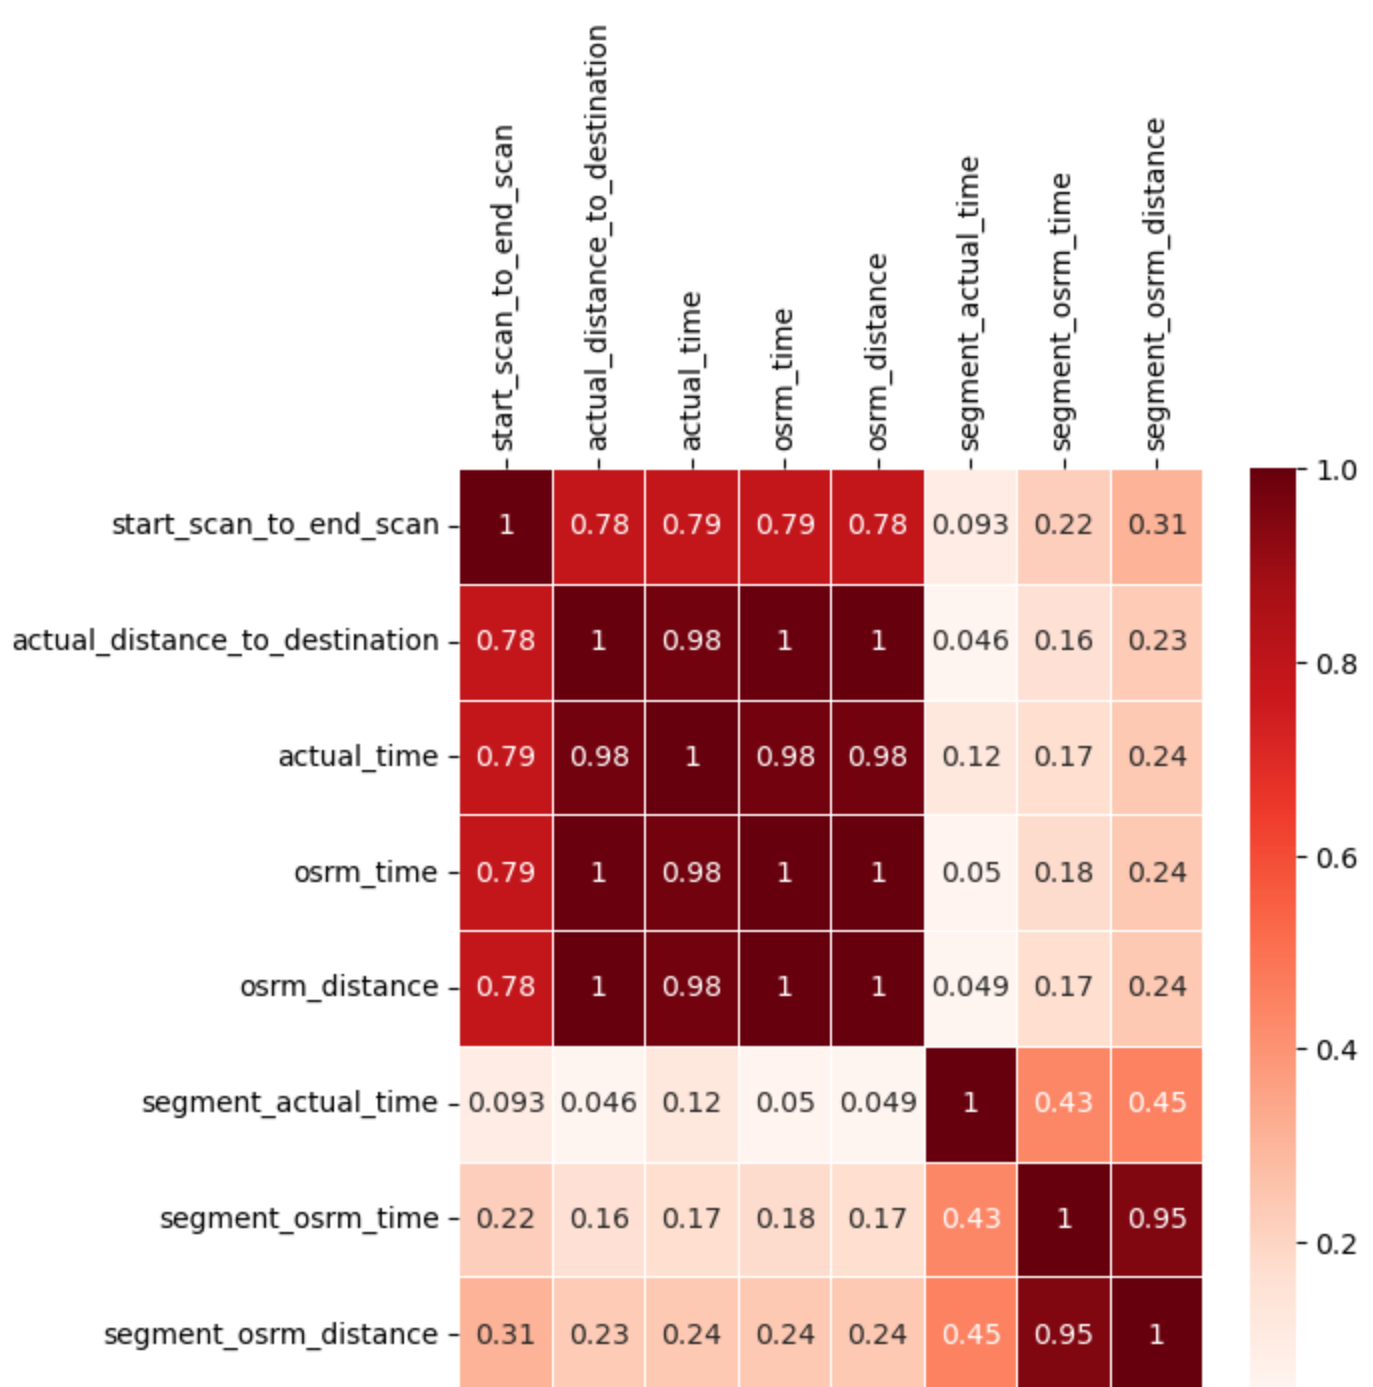


Insight

- The histogram plot of all the **numerical** values show that all the **data is right skewed**
- **72.5%** of the data is **training** data and remaining **27.5%** is **testing** data
- **68.7%** of the delivery is done via **FTL** and remaining **31.3%** through **Carting**

4.3. Multivariate analysis

```
In [14]: fig, ax = plt.subplots(figsize=(6,6))
sns.heatmap(df.select_dtypes(include=np.number).corr(), annot=True, linewidth=0.5, cmap
ax.xaxis.tick_top()
plt.xticks(rotation=90)
plt.show()
```



Insight

- The heatmap clearly shows high correlation between time and distance. This is expected as the delivery time increases with increase in distance
- *actual_distance_to_destination*, *actual_time*, *osrm_time* and *osrm_distance* are highly correlated
- *segment_osrm_time* and *segment_osrm_distance* are highly correlated

4.4. Merging rows

The delivery details of one package is divided into several rows. Creating a unique identifier, called *segment_key*, for different segments of a trip based on the combination of *trip_uuid*, *source_center*, and *destination_center* and then merge the rows of columns *segment_actual_time*, *segment_osrm_distance* and *segment_osrm_time* with same *segment_key* to form new columns *segment_actual_time_sum*, *segment_osrm_distance_sum*, *segment_osrm_time_sum*

```
In [15]: df["segment_key"] = df["trip_uuid"] + '_' + df["source_center"] + '_' + df["destination_center"]
df = df.drop(columns=["source_center", "destination_center"])
segment_columns = ["segment_actual_time", "segment_osrm_distance", "segment_osrm_time"]
for col in segment_columns:
    df[col + "_sum"] = df.groupby("segment_key")[col].cumsum()
```

```
In [16]: df.head(10)
```

Out[16]:

| | data | trip_creation_time | route_schedule_uuid | route_type | trip_uuid | source_name | |
|---|----------|----------------------------|---|------------|--------------------|-------------------------------|----|
| 0 | training | 2018-09-20 02:35:36.476840 | thanos::sroute:eb7bfc78-b351-4c0e-a951-fa3d5c3... | Carting | 153741093647649320 | Anand_VUNagar_DC (Gujarat) | Kh |
| 1 | training | 2018-09-20 02:35:36.476840 | thanos::sroute:eb7bfc78-b351-4c0e-a951-fa3d5c3... | Carting | 153741093647649320 | Anand_VUNagar_DC (Gujarat) | Kh |
| 2 | training | 2018-09-20 02:35:36.476840 | thanos::sroute:eb7bfc78-b351-4c0e-a951-fa3d5c3... | Carting | 153741093647649320 | Anand_VUNagar_DC (Gujarat) | Kh |
| 3 | training | 2018-09-20 02:35:36.476840 | thanos::sroute:eb7bfc78-b351-4c0e-a951-fa3d5c3... | Carting | 153741093647649320 | Anand_VUNagar_DC (Gujarat) | Kh |
| 4 | training | 2018-09-20 02:35:36.476840 | thanos::sroute:eb7bfc78-b351-4c0e-a951-fa3d5c3... | Carting | 153741093647649320 | Anand_VUNagar_DC (Gujarat) | Kh |
| 5 | training | 2018-09-20 02:35:36.476840 | thanos::sroute:eb7bfc78-b351-4c0e-a951-fa3d5c3... | Carting | 153741093647649320 | Khambhat_MotvdDPP_D (Gujarat) | |
| 6 | training | 2018-09-20 02:35:36.476840 | thanos::sroute:eb7bfc78-b351-4c0e-a951-fa3d5c3... | Carting | 153741093647649320 | Khambhat_MotvdDPP_D (Gujarat) | |
| 7 | training | 2018-09-20 02:35:36.476840 | thanos::sroute:eb7bfc78-b351-4c0e-a951-fa3d5c3... | Carting | 153741093647649320 | Khambhat_MotvdDPP_D (Gujarat) | |
| 8 | training | 2018-09-20 02:35:36.476840 | thanos::sroute:eb7bfc78-b351-4c0e-a951-fa3d5c3... | Carting | 153741093647649320 | Khambhat_MotvdDPP_D (Gujarat) | |
| 9 | training | 2018-09-20 02:35:36.476840 | thanos::sroute:eb7bfc78-b351-4c0e-a951-fa3d5c3... | Carting | 153741093647649320 | Khambhat_MotvdDPP_D (Gujarat) | |

10 rows × 21 columns

Grouping the data by *segment_key*, with aggregation defined for each column, and creating a new dataframe *segment*

```
In [17]: segment_dict = {
    'data' : 'first',
    'trip_creation_time' : 'first',
    'route_schedule_uuid' : 'first',
    'route_type' : 'first',
    'trip_uuid' : 'first',
    'source_name' : 'first',
    'destination_name' : 'last',
    'od_start_time' : 'first',
```

```

        'od_end_time' : 'last',
        'start_scan_to_end_scan' : 'first',
        'actual_distance_to_destination' : 'last',
        'actual_time' : 'last',
        'osrm_time' : 'last',
        'osrm_distance' : 'last',
        'segment_actual_time_sum' : 'last',
        'segment_osrm_distance_sum' : 'last',
        'segment_osrm_time_sum' : 'last',
    }

segment_df = df.groupby('segment_key').agg(segment_dict).reset_index()
segment_df = segment_df.sort_values(by=['segment_key', 'od_end_time'], ascending=True).r

```

In [18]: `segment_df.head()`

Out[18]:

| | index | segment_key | data | trip_creation_time | route_schedule_uuid | r |
|---|-------|--|---------------|-------------------------------|---|---|
| 0 | 0 | 153671041653548748_IND209304AAA_IND000000ACB | trip-training | 2018-09-12 00:00:16.535741 | thanos::sroute:d7c989ba-a29b-4a0b-b2f4-288cdc6... | |
| 1 | 1 | 153671041653548748_IND462022AAA_IND209304AAA | trip-training | 2018-09-12 00:00:16.535741 | thanos::sroute:d7c989ba-a29b-4a0b-b2f4-288cdc6... | |
| 2 | 2 | 153671042288605164_IND561203AAB_IND562101AAA | trip-training | 2018-09-12 00:00:22.886430 | thanos::sroute:3a1b0ab2-bb0b-4c53-8c59-eb2a2c0... | |
| 3 | 3 | 153671042288605164_IND572101AAA_IND561203AAB | trip-training | 2018-09-12 00:00:22.886430 | thanos::sroute:3a1b0ab2-bb0b-4c53-8c59-eb2a2c0... | |
| 4 | 4 | 153671043369099517_IND000000ACB_IND160002AAC | trip-training | 2018-09-12 00:00:33.691250 | thanos::sroute:de5e208e-7641-45e6-8100-4d9fb1e... | |

4.5. Feature Engineering

Extracting features from given data

Extracting time taken between *od_start_time* and *od_end_time*

In [19]:

```

segment_df['od_time_diff_hour'] = (segment_df['od_end_time'] - segment_df['od_start_time']) / 3600
segment_df = segment_df.drop(columns=['od_end_time', 'od_start_time'])

```

In [20]: `segment_df.head()`

Out[20]:

| | index | segment_key | data | trip_creation_time | route_schedule_uuid | r |
|---|-------|--|---------------|-------------------------------|---|---|
| 0 | 0 | 153671041653548748_IND209304AAA_IND000000ACB | trip-training | 2018-09-12 00:00:16.535741 | thanos::sroute:d7c989ba-a29b-4a0b-b2f4-288cdc6... | |
| 1 | 1 | 153671041653548748_IND462022AAA_IND209304AAA | trip-training | 2018-09-12 00:00:16.535741 | thanos::sroute:d7c989ba-a29b-4a0b-b2f4-288cdc6... | |
| 2 | 2 | 153671042288605164_IND561203AAB_IND562101AAA | trip-training | 2018-09-12 00:00:22.886430 | thanos::sroute:3a1b0ab2-bb0b-4c53-8c59-eb2a2c0... | |

| | | | | | | |
|---|---|--|----------------|------------|-----------------|---|
| 3 | 3 | 153671042288605164_IND572101AAA_IND561203AAB | trip- training | 2018-09-12 | 00:00:22.886430 | thanos::sroute:3a1b0ab2-bb0b-4c53-8c59-eb2a2c0... |
| 4 | 4 | 153671043369099517_IND000000ACB_IND160002AAC | trip- training | 2018-09-12 | 00:00:33.691250 | thanos::sroute:de5e208e-7641-45e6-8100-4d9fb1e... |

Extracting city, place, code and state from *source_name* and *destination_name*

```
In [21]: segment_df['source_state'] = segment_df['source_name'].str.extract(r'\((.*?)\)')
segment_df['source_data'] = segment_df['source_name'].str.extract(r'^(.*)\()')
segment_df['source_data'] = segment_df['source_data'].str.strip()

segment_df['destination_state'] = segment_df['destination_name'].str.extract(r'\((.*?)\)')
segment_df['destination_data'] = segment_df['destination_name'].str.extract(r'^(.*)\()')
segment_df['destination_data'] = segment_df['destination_data'].str.strip()
```

```
In [22]: def extract_city_place_code(name):
    parts = name.split('_')
    num_of_parts = len(parts)
    if(num_of_parts == 3):
        city = parts[0]
        place = parts[1]
        code = parts[2]
    elif(num_of_parts == 2):
        city = parts[0]
        place = parts[1]
        code = 'none'
    else:
        city = parts[0]
        place = city
        code = 'none'

    if city == 'Bangalore' or city == 'HBR Layout PC' or city == 'BLR':
        city = 'Bengaluru'
    elif city == 'Mumbai Hub' or city == 'BOM':
        city = 'Mumbai'
    elif city == 'Del':
        city = 'Delhi'
    elif city == 'PNQ Pashan DPC' or city == 'PNQ Vadgaon Sheri DPC':
        city = 'Pune'
    elif city == 'MAA':
        city = 'Chennai'
    elif city == 'FBD':
        city = 'Faridabad'
    elif city == 'CCU':
        city = 'Kolkata'
    elif city == 'AMD':
        city = 'Ahmedabad'
    elif city == 'FBD':
        city = 'Faridabad'
    elif city == 'GGN':
        city = 'Gurgaon'
    elif city == 'GZB':
        city = 'Ghaziabad'

    return [city, place, code]
```

```
In [23]: extracted_df = segment_df['source_data'].apply(lambda x: extract_city_place_code(x))
segment_df[['source_city', 'source_place', 'source_code']] = pd.DataFrame(extracted_df.tolist())
extracted_df = segment_df['destination_data'].apply(lambda x: extract_city_place_code(x))
```

```
segment_df[['destination_city', 'destination_place', 'destination_code']] = pd.DataFrame(e
segment_df = segment_df.drop(columns=['source_name', 'source_data', 'destination_name',
segment_df.head()
```

Out[23]:

| | index | segment_key | data | trip_creation_time | route_schedule_uuid | r |
|---|-------|--|---------------|-------------------------------|---|---|
| 0 | 0 | 153671041653548748_IND209304AAA_IND000000ACB | trip-training | 2018-09-12 00:00:16.535741 | thanos::sroute:d7c989ba-a29b-4a0b-b2f4-288cdc6... | |
| 1 | 1 | 153671041653548748_IND462022AAA_IND209304AAA | trip-training | 2018-09-12 00:00:16.535741 | thanos::sroute:d7c989ba-a29b-4a0b-b2f4-288cdc6... | |
| 2 | 2 | 153671042288605164_IND561203AAB_IND562101AAA | trip-training | 2018-09-12 00:00:22.886430 | thanos::sroute:3a1b0ab2-bb0b-4c53-8c59-eb2a2c0... | |
| 3 | 3 | 153671042288605164_IND572101AAA_IND561203AAB | trip-training | 2018-09-12 00:00:22.886430 | thanos::sroute:3a1b0ab2-bb0b-4c53-8c59-eb2a2c0... | |
| 4 | 4 | 153671043369099517_IND000000ACB_IND160002AAC | trip-training | 2018-09-12 00:00:33.691250 | thanos::sroute:de5e208e-7641-45e6-8100-4d9fb1e... | |

5 rows × 24 columns

In [24]:

```
segment_df['trip_creation_year'] = segment_df['trip_creation_time'].dt.year
segment_df['trip_creation_month'] = segment_df['trip_creation_time'].dt.month
segment_df['trip_creation_day'] = segment_df['trip_creation_time'].dt.day
segment_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26222 entries, 0 to 26221
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   index                                26222 non-null  int64
1   segment_key                          26222 non-null  object
2   data                                26222 non-null  category
3   trip_creation_time                  26222 non-null  datetime64[ns]
4   route_schedule_uuid                 26222 non-null  object
5   route_type                          26222 non-null  category
6   trip_uuid                           26222 non-null  object
7   start_scan_to_end_scan              26222 non-null  float64
8   actual_distance_to_destination      26222 non-null  float64
9   actual_time                         26222 non-null  float64
10  osrm_time                           26222 non-null  float64
11  osrm_distance                       26222 non-null  float64
12  segment_actual_time_sum             26222 non-null  float64
13  segment_osrm_distance_sum           26222 non-null  float64
14  segment_osrm_time_sum               26222 non-null  float64
15  od_time_diff_hour                   26222 non-null  float64
16  source_state                        26222 non-null  object
17  destination_state                   26222 non-null  object
18  source_city                         26222 non-null  object
19  source_place                        26222 non-null  object
20  source_code                         26222 non-null  object
21  destination_city                    26222 non-null  object
22  destination_place                   26222 non-null  object
23  destination_code                    26222 non-null  object
24  trip_creation_year                  26222 non-null  int32
25  trip_creation_month                 26222 non-null  int32
26  trip_creation_day                   26222 non-null  int32
```


dtypes: category(2), datetime64[ns](1), float64(9), int32(3), int64(1), object(11)
memory usage: 4.8+ MB

4.6. In-depth analysis

4.6.1. Grouping and aggregating at trip-level

Group the *segment_df* by *trip_uuid*

```
In [25]: trip_dict = {
    'segment_key' : 'first',
    'data' : 'first',
    'trip_creation_time' : 'first',
    'route_schedule_uuid' : 'first',
    'route_type' : 'first',
    'start_scan_to_end_scan' : 'sum',
    'actual_distance_to_destination' : 'sum',
    'actual_time' : 'sum',
    'osrm_time' : 'sum',
    'osrm_distance' : 'sum',
    'segment_actual_time_sum' : 'sum',
    'segment_osrm_distance_sum' : 'sum',
    'segment_osrm_time_sum' : 'sum',
    'od_time_diff_hour' : 'sum',
    'source_state' : 'first',
    'destination_state' : 'last',
    'source_city' : 'first',
    'source_place' : 'first',
    'source_code' : 'first',
    'destination_city' : 'last',
    'destination_place' : 'last',
    'destination_code' : 'last',
}
trip_df = segment_df.groupby('trip_uuid').agg(trip_dict).reset_index()
```

```
In [26]: trip_df.head()
```


```
Out[26]:
```

| | trip_uuid | segment_key | data | trip_creation_time | route_scl |
|---|-----------------------------|---|----------|-------------------------------|-----------------------|
| 0 | trip- 153671041653548748 | trip- 153671041653548748_IND209304AAA_IND000000ACB | training | 2018-09-12 00:00:16.535741 | thanos::srout a29t |
| 1 | trip- 153671042288605164 | trip- 153671042288605164_IND561203AAB_IND562101AAA | training | 2018-09-12 00:00:22.886430 | thanos::srout bb0t |
| 2 | trip- 153671043369099517 | trip- 153671043369099517_IND000000ACB_IND160002AAC | training | 2018-09-12 00:00:33.691250 | thanos::srout 7641 |
| 3 | trip- 153671046011330457 | trip- 153671046011330457_IND400072AAB_IND401104AAA | training | 2018-09-12 00:01:00.113710 | thanos::srout a679 |
| 4 | trip- 153671052974046625 | trip- 153671052974046625_IND583101AAA_IND583201AAA | training | 2018-09-12 00:02:09.740725 | thanos::srout 65eC |

5 rows × 23 columns


4.6.2. Outlier Detection & Treatment

```
In [27]: trip_numerical_columns = ['start_scan_to_end_scan', 'actual_distance_to_destination',
                                   'actual_time', 'osrm_time', 'osrm_distance', 'segment_actual_t',
                                   'segment_osrm_distance_sum', 'segment_osrm_time_sum', 'od_time',
                                   'od_time_diff_hour']
trip_df[trip_numerical_columns].boxplot(figsize=(25,12))
plt.show()
```



```
In [28]: Q1 = trip_df[trip_numerical_columns].quantile(0.25)
Q3 = trip_df[trip_numerical_columns].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5*IQR
higher_bound = Q3 + 1.5*IQR
trip_df = trip_df[-((trip_df[trip_numerical_columns] < lower_bound) | (trip_df[trip_nu
trip_df = trip_df.reset_index(drop=True))
```

```
In [29]: trip_df[trip_numerical_columns].boxplot(figsize=(25,12))
plt.show()
```



4.6.3. Perform one-hot encoding on categorical features

route_type is the only categorical feature

```
In [30]: ohe_df = pd.get_dummies(trip_df['route_type'], dtype='int', prefix='route_type')
trip_df = pd.concat([trip_df, ohe_df], axis=1)
trip_df = trip_df.drop(columns='route_type')
```

```
In [31]: trip_df.head()
```

```
Out[31]:
```

| | trip_uuid | segment_key | data | trip_creation_time | route_scl |
|---|-------------------------|---|----------|----------------------------|----------------------|
| 0 | trip-153671042288605164 | trip-153671042288605164_IND561203AAB_IND562101AAA | training | 2018-09-12 00:00:22.886430 | thanos::srou bb0t |
| 1 | trip-153671046011330457 | trip-153671046011330457_IND400072AAB_IND401104AAA | training | 2018-09-12 00:01:00.113710 | thanos::srou a679 |
| 2 | trip-153671052974046625 | trip-153671052974046625_IND583101AAA_IND583201AAA | training | 2018-09-12 00:02:09.740725 | thanos::srou 65e6 |
| 3 | trip-153671055416136166 | trip-153671055416136166_IND600056AAA_IND602105AAB | training | 2018-09-12 00:02:34.161600 | thanos::srou d0a2 |
| 4 | trip-153671066201138152 | trip-153671066201138152_IND600044AAD_IND600048AAA | training | 2018-09-12 00:04:22.011653 | thanos::srou 846e |

5 rows × 24 columns

```
In [32]: print('Number of Carting route is ', trip_df[trip_df['route_type_Carting'] == 1]['route_type_Carting'].value_counts())
print('Number of FTL route is ', trip_df[trip_df['route_type_FTL'] == 1]['route_type_FTL'].value_counts())
```

Number of Carting route is 8812
Number of FTL route is 3911


```
In [33]: plt.figure(figsize=(8,8))
data = trip_df["source_state"]
ax=sns.countplot(y = data, order=data.value_counts().index)
ax.bar_label(ax.containers[0])
plt.show()


plt.figure(figsize=(8,8))
data = trip_df["source_city"]
ax=sns.countplot(y = data, order=data.value_counts()[:20].index)
ax.bar_label(ax.containers[0])
plt.show()


plt.figure(figsize=(8,8))
data = trip_df["destination_state"]
ax=sns.countplot(y = data, order=data.value_counts().index)
ax.bar_label(ax.containers[0])
plt.show()


plt.figure(figsize=(8,8))
data = trip_df["destination_city"]
ax=sns.countplot(y = data, order=data.value_counts()[:20].index)
ax.bar_label(ax.containers[0])
plt.show()
```

source_state










```
In [34]: trip_df["corridor"] = trip_df["source_city"] + '_' + trip_df["destination_city"]
plt.figure(figsize=(8,4))
ax=sns.countplot(y = trip_df["corridor"], order=trip_df["corridor"].value_counts()[:20])
ax.bar_label(ax.containers[0])
plt.show()
```



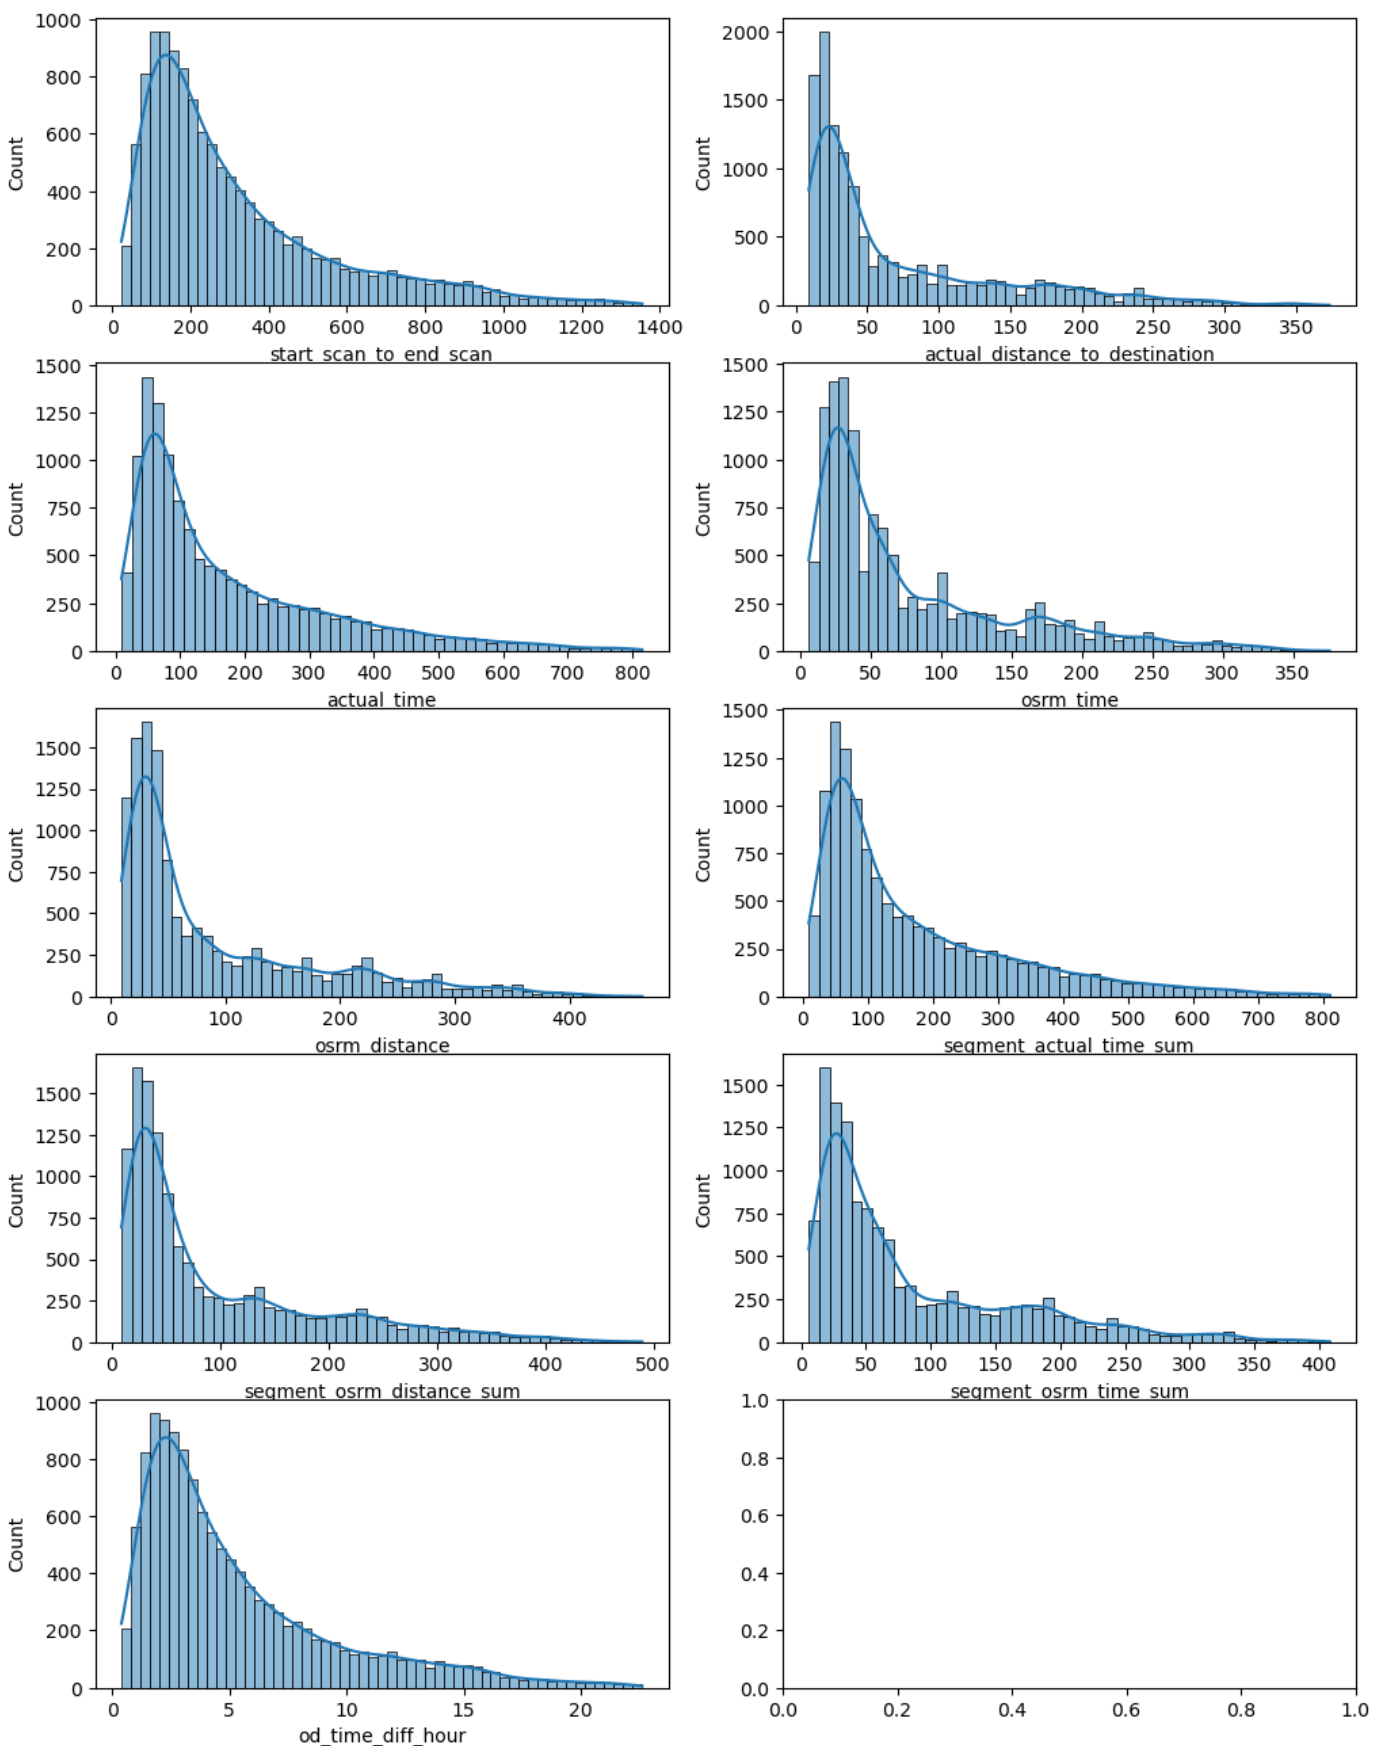
```
In [35]: Mumbai_Bhiwandi_df = trip_df[((trip_df["corridor"] == "Bhiwandi_Mumbai") | (trip_df["corridor"] == "Bhiwandi_Mumbai"))]
print('Avg time: ', Mumbai_Bhiwandi_df['actual_time'].mean())
print('Avg distance: ', Mumbai_Bhiwandi_df['actual_distance_to_destination'].mean())
```

Avg time: 81.8186968838527

Avg distance: 22.218624868058914

4.6.4. Normalize/ Standardize the numerical features using MinMaxScaler or StandardScaler

```
In [36]: fig, ax = plt.subplots(nrows=5, ncols=2, figsize = (12, 16))
sns.histplot(data=trip_df, x = "start_scan_to_end_scan", kde=True, ax=ax[0,0])
sns.histplot(data=trip_df, x = "actual_distance_to_destination", kde=True, ax=ax[0,1])
sns.histplot(data=trip_df, x = "actual_time", kde=True, ax=ax[1,0])
sns.histplot(data=trip_df, x = "osrm_time", kde=True, ax=ax[1,1])
sns.histplot(data=trip_df, x = "osrm_distance", kde=True, ax=ax[2,0])
sns.histplot(data=trip_df, x = "segment_actual_time_sum", kde=True, ax=ax[2,1])
sns.histplot(data=trip_df, x = "segment_osrm_distance_sum", kde=True, ax=ax[3,0])
sns.histplot(data=trip_df, x = "segment_osrm_time_sum", kde=True, ax=ax[3,1])
sns.histplot(data=trip_df, x = "od_time_diff_hour", kde=True, ax=ax[4,0])
plt.show()
```

Insight

- None of the data is gaussian, so we will use MinMaxScaler

```
In [37]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
```

```
scaler.fit(trip_df[trip_numerical_columns])
trip_df[trip_numerical_columns] = scaler.transform(trip_df[trip_numerical_columns])
```

```
In [38]: trip_df.describe()
```


| | trip_creation_time | start_scan_to_end_scan | actual_distance_to_destination | actual_time | osrm_time | osrm |
|--------------|----------------------------------|------------------------|--------------------------------|--------------|--------------|-------|
| count | 12723 | 12723.000000 | 12723.000000 | 12723.000000 | 12723.000000 | 12723 |
| mean | 2018-09-22 13:16:08.771620608 | 0.223107 | 0.173734 | 0.208998 | 0.195785 | |
| min | 2018-09-12 00:00:22.886430 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 25% | 2018-09-17 03:12:27.545116928 | 0.084835 | 0.034006 | 0.064516 | 0.056757 | |
| 50% | 2018-09-22 04:23:52.568071936 | 0.157658 | 0.081009 | 0.130273 | 0.118919 | |
| 75% | 2018-09-27 20:46:53.577142016 | 0.300300 | 0.254284 | 0.300248 | 0.278378 | |
| max | 2018-10-03 23:59:42.701692 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | |
| std | NaN | 0.191859 | 0.197757 | 0.196217 | 0.195496 | |

4.7. Hypothesis Testing

4.7.1. Are aggregated *actual_time* and aggregated *osrm_time* similar?

H0 : *actual_time* and *osrm_time* are similar \ H1 : *actual_time* and *osrm_time* are different

```
In [39]: plt.figure(figsize=(10,6))
sns.histplot(trip_df['actual_time'], color='green')
sns.histplot(trip_df['osrm_time'], color='blue')
plt.legend(['actual_time', 'osrm_time'])
plt.show()
```



This is a 2 sample continuous skewed data, so we will use Mann-Whitney U Test

```
In [40]: statistic, pvalue = sps.mannwhitneyu(trip_df['actual_time'], trip_df['osrm_time'], alter
print('p-value', pvalue)
if pvalue < 0.1:
    print('The samples are not similar')
else:
    print('The samples are similar ')
```

```
p-value 1.3094485692382313e-20
The samples are not similar
```


Insight

- *actual_time* and *osrm_time* are different

4.7.2. Are aggregated *actual_time* and aggregated *segment_actual_time* similar?

H_0 : *actual_time* and *segment_actual_time* are similar \ H_1 : *actual_time* and *segment_actual_time* are different

```
In [41]: plt.figure(figsize=(10,6))
sns.histplot(trip_df['actual_time'], color='green')
sns.histplot(trip_df['segment_actual_time_sum'], color='blue')
plt.legend(['actual_time', 'segment_actual_time_sum'])
plt.show()
```



This is a 2 sample continuous skewed data, so we will use Mann-Whitney U Test

```
In [42]: statistic, pvalue = sps.mannwhitneyu(trip_df['actual_time'], trip_df['segment_actual_time'])
print('p-value', pvalue)
if pvalue < 0.1:
    print('The samples are not similar')
else:
    print('The samples are similar ')
```


p-value 0.7167057478572094

The samples are similar

4.7.3. Are aggregated *osrm_distance* and aggregated *segment_osrm_distance* similar?

H0 : *osrm_distance* and *segment_osrm_distance* are similar \ H1 : *osrm_distance* and *segment_osrm_distance* are different

```
In [43]: plt.figure(figsize=(10,6))
sns.histplot(trip_df['osrm_distance'], color='green')
sns.histplot(trip_df['segment_osrm_distance_sum'], color='blue')
plt.legend(['osrm_distance', 'segment_osrm_distance_sum'])
plt.show()
```



This is a 2 sample continuous skewed data, so we will use Mann-Whitney U Test


```
In [44]: statistic, pvalue = sps.mannwhitneyu(trip_df['osrm_distance'], trip_df['segment_osrm_dis
print('p-value', pvalue)
if pvalue < 0.1:
    print('The samples are not similar')
else:
    print('The samples are similar ')
```

```
p-value 0.05751040543671224
The samples are not similar
```

4.7.4. Are aggregated *osrm_time* and aggregated *segment_osrm_time* similar?

H0 : *osrm_time* and *segment_osrm_time* are similar \ H1 : *osrm_time* and *segment_osrm_time* are different

```
In [45]: plt.figure(figsize=(10,6))
sns.histplot(trip_df['osrm_time'], color='green')
sns.histplot(trip_df['segment_osrm_time_sum'], color='blue')
plt.legend(['osrm_time', 'segment_osrm_time_sum'])
plt.show()
```



This is a 2 sample continuous skewed data, so we will use Mann-Whitney U Test

```
In [46]: statistic, pvalue = sps.mannwhitneyu(trip_df['osrm_time'], trip_df['segment_osrm_time_sum'])
print('p-value', pvalue)
if pvalue < 0.1:
    print('The samples are not similar')
else:
    print('The samples are similar')
```

p-value 0.8230933178296898
The samples are similar

5. Business Insights

- The most common route type is **Carting**
- The top 3 **source states** are **Maharashtra, Karnataka and Haryana**
- The top 3 **source cities** are **Bengaluru, Mumbai and Gurgaon**
- The top 3 **destination states** are **Maharashtra, Karnataka and Haryana**
- The top 3 **destination cities** are **Bengaluru, Mumbai and Gurgaon**
- Most of the packages are sent and received within Bengaluru, Mumbai and Chennai but the **most bussiest corridor is Bhiwandi-Mumbai**
- Aggregated **actual_time** and aggregated **osrm_time** are **not similar**
- Aggregated **actual_time** and aggregated **segment_actual_time** are **similar**
- Aggregated **osrm_distance** and aggregated **segment_osrm_distance** are **not similar**
- Aggregated **osrm_time** and aggregated **segment_osrm_time** are **similar**

6. Recommendation

- The company should advertise more on route type FTL saying it is faster mode of delivery. This way FTL can be suggested to atleast large organization.
- Cities Bengaluru(Karnataka), Mumbai(Maharashtra) and Gurgaon(Haryana) send and receive the majority of the deliveries. The company should keep the customers of these cities satisfied with the better and faster services. This involves improving the OSRM engine to make better delivery time predictions.

In []: